

Introduction to Machine Learning

Lab 1 Block 2

Anton Persson, Emil Klasson Svensson, Mattias Karlsson, Rasmus Holm

2016-11-21

Contents

Assignment 1	2
2	2
3	2
Assignment 2	4
1	4
2	4
3	5
4	6
5	9
6	9
Appendix	11
Code for Assignment 1	11
Code for Assignment 2	11
Contributions	13

Assignment 1

2

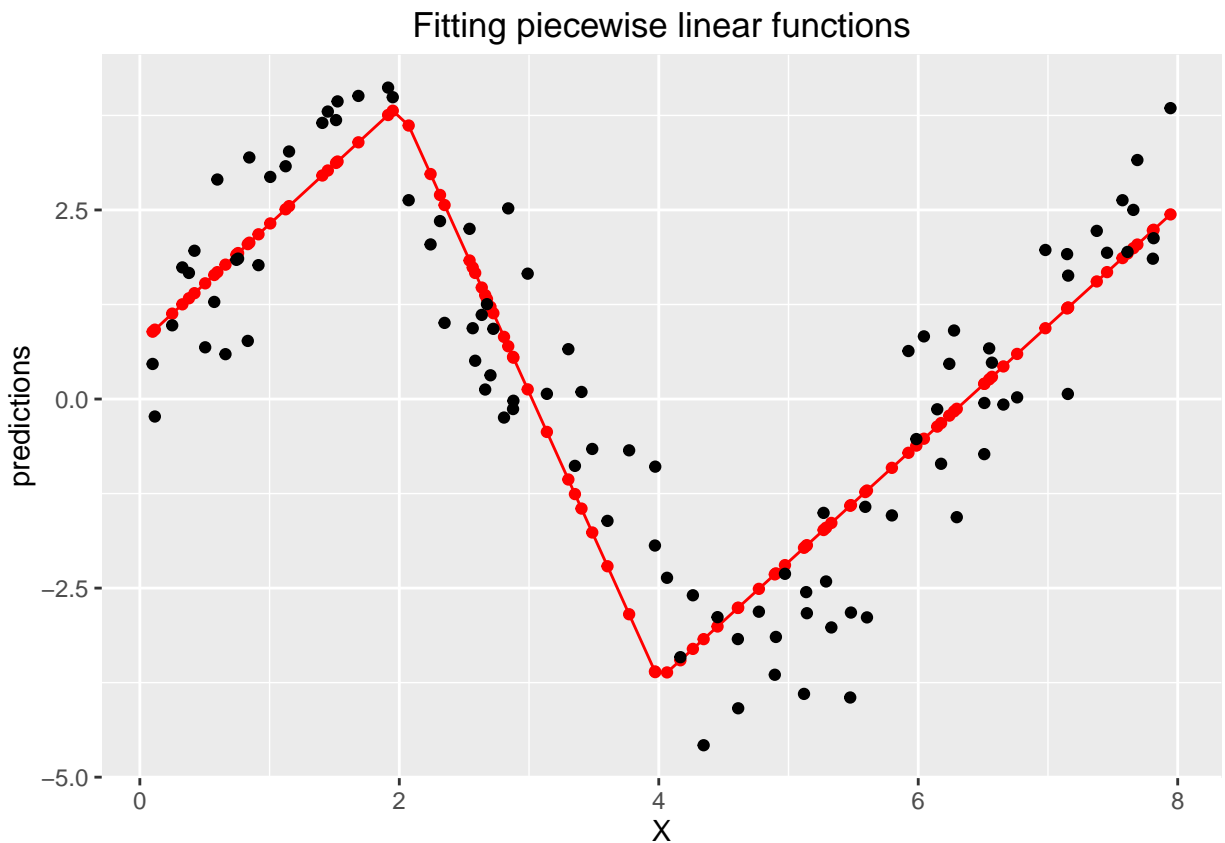


Figure 1: Piecewise Linear Model

The quality of the fit in figure 1 seems to be reasonable good, considering that only two knots was used at arbitrary placements. Would the placement been chosen more wisely the fit would have been even better, for instance at around $x = 1.9$ and $x = 4.8$. The function is continuous piecewise linear with derivatives of order one everywhere except at the knots/end points.

3

In this exercise we used the `smooth.spline()` instead which returns a continuous cubic spline function and the resulting fit can be seen in figure 2

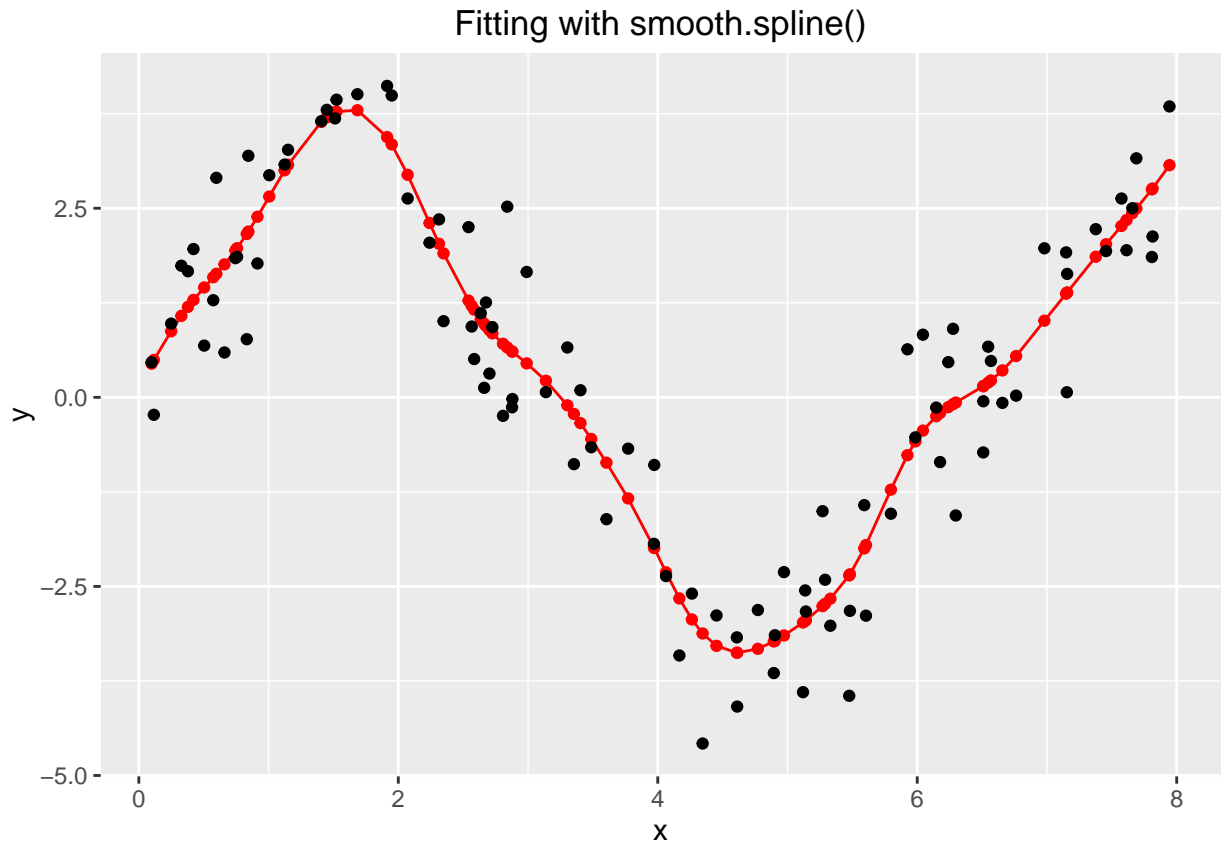


Figure 2: Cubic Smooth Spline

The fit is definitively better than the piecewise linear, it is more precise and catches the trends in the data. However, the cubic spline runs the risk of overfitting the data and the piecewise linear function would perhaps be a much better choice with a more thought-out placement of the knots. The data also look like it have been generated by three different linear functions which the piecewise linear function could match reasonably well.

Assignment 2

In the following exercises I have used the influenza data set that contains weekly data on the mortality and the number of laboratory-confirmed cases of influenza in Sweden.

1




Figure 3: Influenza and mortality cases as time series between 1995 and 2004

From figure 3 we can see that spikes in influenza also resulted in increased mortality cases over the complete timeline. This indicates that mortality cases are positively correlated with influenza.

2

In this exercise I assume $Mortality \sim \mathcal{N}(\mu, \sigma^2)$ so the probabilistic model becomes

$$g(\mathbf{E}[Mortality \mid Year, Week]) = \alpha + \beta_1 Year + f(Week),$$

where $g(\mu) = \mu$ is the link function and f is a spline with an estimated degrees of freedom of 8.487.




Figure 4: Fitted model to the actual data.

Figure 4 shows that the generalized additive model (GAM) is a decent fit to the data. The true data is very chaotic which the model does not capture, i.e. do not overfit to the noise, but it does underfit the peaks so overall I would say the model underfits the data.

We can also see that the mortality outbreak occurs in the beginning of each year and then decreases until winter is coming again, so the mortality is positively correlated with influenza cases which in turn are increased by the colder climate in the late/early months of each year.

```
#>
#> Family: gaussian
#> Link function: identity
#>
#> Formula:
#> Mortality ~ Year + s(Week)
#>
#> Parametric coefficients:
#>             Estimate Std. Error t value Pr(>|t|)
#> (Intercept) -652.058   3448.379  -0.189    0.85
#> Year          1.219     1.725    0.706    0.48
#>
#> Approximate significance of smooth terms:
#>             edf Ref.df    F p-value
```

```
#> s(Week) 8.587 8.951 100.6 <2e-16 ***
#> ---
#> Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
#>
#> R-sq.(adj) = 0.661   Deviance explained = 66.8%
#> GCV = 9014.6   Scale est. = 8806.7    n = 459
```

From the model summary above we can see that the spline component is statistical significant while the intercept and year variable are not. This is not surprising given that the data is non-linear from figure 4.




Figure 5: The week spline function.

The fitted *week* spline function can be seen in figure 5 with the pointwise standard errors included. It shows that the value is high in the early/late weeks and it is at its lowest point in the middle of the year, i.e. the summer months. This correspond to the same analysis as previously that mortality cases increase in the early/late months each year.

4

We know that the smoothing penalty factor, λ , and degrees of freedom for a spline have the following relationship.

$$\text{df}_\lambda = \sum_{k=1}^N \frac{1}{1 + \lambda d_k}$$

implicates that as λ increases the degrees of freedom decreases and vice versa. Figure 6 shows how the fit changes with varying penalty factors (sp parameter) with the maximum degrees of freedom set to a fixed value of 51 (k parameter).

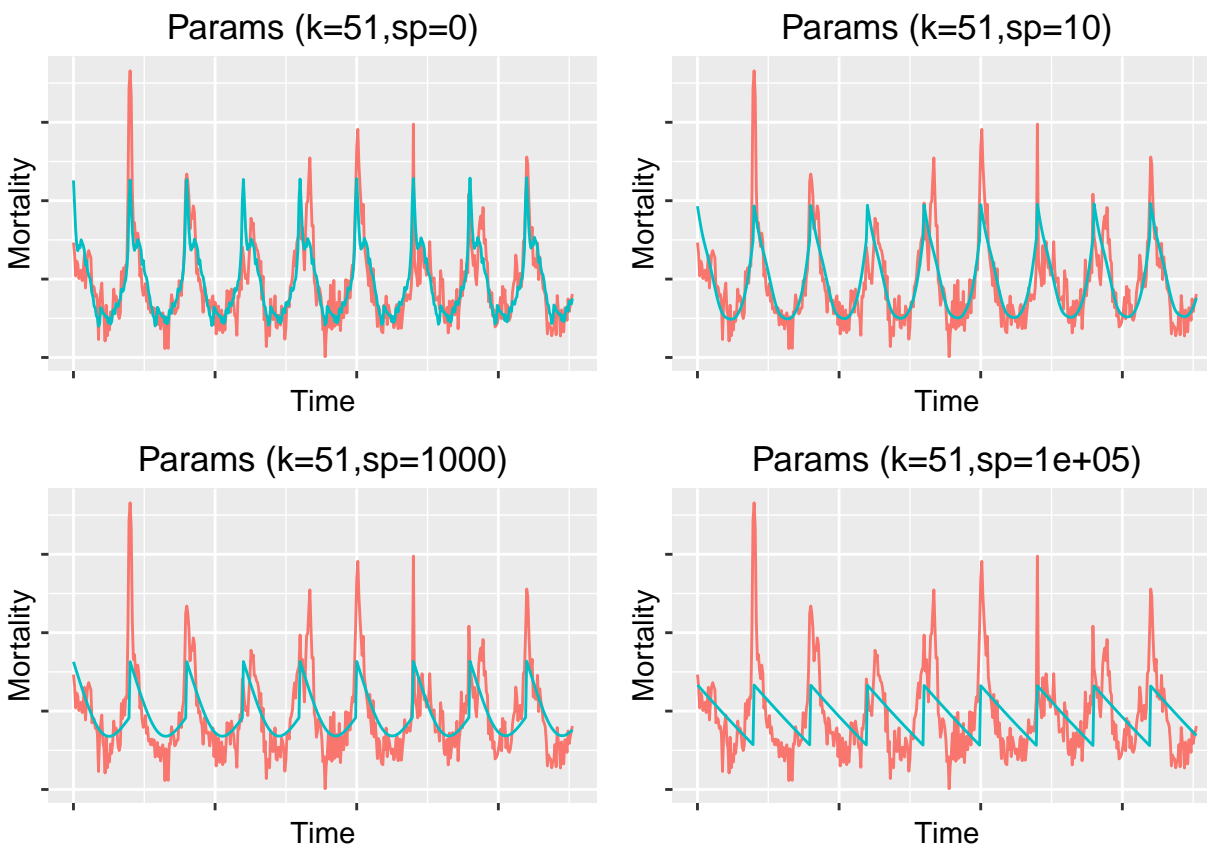


Figure 6: Shows how the estimation changes with varying the penalty factor of the GAM.

As the penalty factor increases the fit becomes less flexible due to a decrease in the degrees of freedom which results in a fit that does not match the data particular well and this follows directly from the theory above. A penalty factor of 0 or 10 are decent in this case but the model is still underfitting the real underlying data which means that the model is not complex enough. The maximum degrees of freedom possible is used when the penalty factor is set to 0 and therefore indicates the result that the *week* feature cannot explain the mortality completely so a better model would need to incorporate more knowledge.




Figure 7: Deviance versus penalty factor.

Figure 7 shows that as the penalty factor increases so does the estimated deviance of the model which is not desirable.

5




Figure 8: Residuals from the GAM in exercise 2.

The residuals are usually negative when *influenza* is at 0 and large positive spikes when the *influenza* is above 0 as shown in figure 8. This indicates that the model underestimates during the early/late months and usually overestimates the other months.

6

In this new model I have modelled *mortality* as an additive function of the spline functions of *year*, *week*, and *influenza*.

```
#>
#> Family: gaussian
#> Link function: identity
#>
#> Formula:
#> Mortality ~ s(Year, k = length(unique(data$Year)) - 1) + s(Week,
#>   k = length(unique(data$Week)) - 1) + s(Influenza, k = length(unique(data$Influenza)))
#>
#> Parametric coefficients:
#>             Estimate Std. Error t value Pr(>|t|)
#> (Intercept)  1783.77      3.28   543.9  <2e-16 ***
#> ---
```

```

#> Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
#>
#> Approximate significance of smooth terms:
#>              edf Ref.df      F p-value
#> s(Year)         3.913  4.756  1.179   0.292
#> s(Week)        13.886 17.271 20.271 <2e-16 ***
#> s(Influenza)    59.204 65.527  5.336 <2e-16 ***
#> ---
#> Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
#>
#> Rank: 132/142
#> R-sq.(adj) =   0.81   Deviance explained = 84.2%
#> GCV = 5947.8   Scale est. = 4937       n = 459

```

The model summary above shows that the *influenza* is statistical significant in explaining the mortality and so is the *week* variable like before.




Figure 9: Fitted model to the actual data.

The new model is fitting the data much better than before as shown in figure 9 and from the summary we can conclude that the *week* and *influenza* variables explain most of the mortality cases, roughly 84% of the deviance.

Appendix

Code for Assignment 1

```
myspline<-function(X, Y, knots){  
  
  X<-as.data.frame(cbind(Intercept=rep(1, length(X)), X)) #Lägger till intercept  
  
  for (i in 1:length(knots)){  
  
    X[,i+2]<- pmax(cube$x-knots[i], 0)  
  }  
  all_data<-data.frame(y=Y, X)  
  
  lin_mod<-lm(y~.-Intercept, data=all_data)  
  
  predictions<-predict(lin_mod)  
  
  library(ggplot2)  
  linear_plot<-ggplot(data=all_data)+geom_line(aes(x=X, y=predictions), col="red")+  
    geom_point(aes(x=X, y=predictions), col="red")+geom_point(aes(y=all_data$y, x=all_data$X))+  
    ggtitle("Fitting piecewise linear functions")  
  
  plot(linear_plot)  
  
  #return(summary(lin_mod)) Don't need the summary in the report so I delete that.  
}  
  
cube <- read.csv2("../data/cube.csv", header=TRUE, sep=";")  
myspline(X=cube$x, Y=cube$y, knots=c(2,4))  
  
X<-cube$x  
Y<-cube$y  
  
spline_predictions<-predict(smooth.spline(x=X, y=Y))  
  
spline_predictions<-as.data.frame(cbind(x=as.vector(spline_predictions$x),y=as.vector(spline_predictions$y)))  
  
spline_plot<-ggplot()+geom_point(data=spline_predictions, aes(x=x, y=y), col="red")+  
  geom_line(data=spline_predictions, aes(x=x, y=y), col="red")+  
  geom_point(aes(y=cube$y, x=cube$x))+ggtitle("Fitting with smooth.spline()")  
  
spline_plot
```

Code for Assignment 2

```
library(ggplot2)  
library(readxl)  
library(reshape2)  
library(mgcv)  
library(grid)
```

```

library(gridExtra)

data <- read_excel("../data/Influenza.xlsx")
plot_data <- melt(data[, c("Time", "Mortality", "Influenza")], id="Time")
ggplot(plot_data) +
  geom_line(aes(x=Time, y=value, color=variable)) +
  scale_x_discrete(limit=data$Year)
gamfit <- gam(Mortality ~ Year + s(Week), family=gaussian, data=data, method="GCV.Cp")

yhat <- predict(gamfit, data)

plot_data <- data.frame(Time=data$Time, Actual=data$Mortality, Estimate=as.numeric(yhat))
plot_data <- melt(plot_data, id="Time", value.name="Mortality", variable.name="Data")

ggplot(plot_data) +
  geom_line(aes(x=Time, y=Mortality, color=Data)) +
  scale_x_discrete(limit=data$Year)
summary(gamfit)
plot(gamfit)

k <- length(unique(data$Week)) - 1
penalty_values <- c(0, 10, 1000, 100000)
deviance <- rep(0, length(penalty_values))

plots <- list()

for (i in 1:length(penalty_values)) {
  fit <- gam(Mortality ~ Year + s(Week, k=k, sp=penalty_values[i]),
    family=gaussian, data=data, method="GCV.Cp")

  deviance[i] <- deviance(fit)

  title <- paste("Params (k=", k, ", sp=", penalty_values[i], ")", sep="")

  plot_data <- data.frame(Time=data$Time, Actual=data$Mortality, Estimate=fitted(fit))
  plot_data <- melt(plot_data, id="Time", value.name="Mortality", variable.name="Data")

  plots[[i]] <- ggplot(plot_data) +
    geom_line(aes(x=Time, y=Mortality, color=Data), show.legend=FALSE) +
    ggtitle(title) +
    theme(axis.text=element_blank(),
      plot.title=element_text(hjust=0.5))
}

do.call(grid.arrange, c(plots, list(ncol=2)))
plot_data <- data.frame(Penalty=penalty_values, Deviance=deviance)
ggplot(plot_data) +
  geom_line(aes(x=log(Penalty), y=log(Deviance)))

gamfit <- gam(Mortality ~ Year + s(Week), family=gaussian, data=data, method="GCV.Cp")
residuals <- resid(gamfit)
plot_data <- data.frame(Time=data$Time, Residuals=residuals, Influenza=data$Influenza)
plot_data <- melt(plot_data, id="Time", value.name="Value", variable.name="Data")

```

```

ggplot(plot_data) +
  geom_line(aes(x=Time, y=Value, color=Data)) +
  scale_x_discrete(limit=data$Year)

gamfit <- gam(Mortality ~ s(Year, k=length(unique(data$Year)) - 1) +
              s(Week, k=length(unique(data$Week)) - 1) +
              s(Influenza, k=length(unique(data$Influenza))),
              data=data)
summary(gamfit)
plot_data <- data.frame(Time=data$Time, Actual=data$Mortality, Estimate=fitted(gamfit))
plot_data <- melt(plot_data, id="Time", value.name="Value", variable.name="Data")

ggplot(plot_data) +
  geom_line(aes(x=Time, y=Value, color=Data)) +
  scale_x_discrete(limit=data$Year)

```

Contributions

We divided the work into two parts and discussed/compiled the results in pairs. Then we all discussed our findings together as a whole group and checked that everyone had similar/understood the results.