

Introduction to Machine Learning

Lab 3

Rasmus Holm

2016-11-18

Contents

Assignment 1	2
1	2
2	3
3	4
4	5
Assignment 2	6
1	6
2	6
3	7
4	8
5	8
Appendix	10
Code for Assignment 1	10
Code for Assignment 2	11

Assignment 1

1



2



3



4



Assignment 2

1

2


```
#> $confusion_matrix
#>      true
#> pred  bad good
#>   bad   61  20
#>   good  86 333
#>
#> $classification_rate
#> [1] 0.788

#> $confusion_matrix
#>      true
#> pred  bad good
#>   bad   34  21
#>   good  41 154
#>
#> $classification_rate
#> [1] 0.752

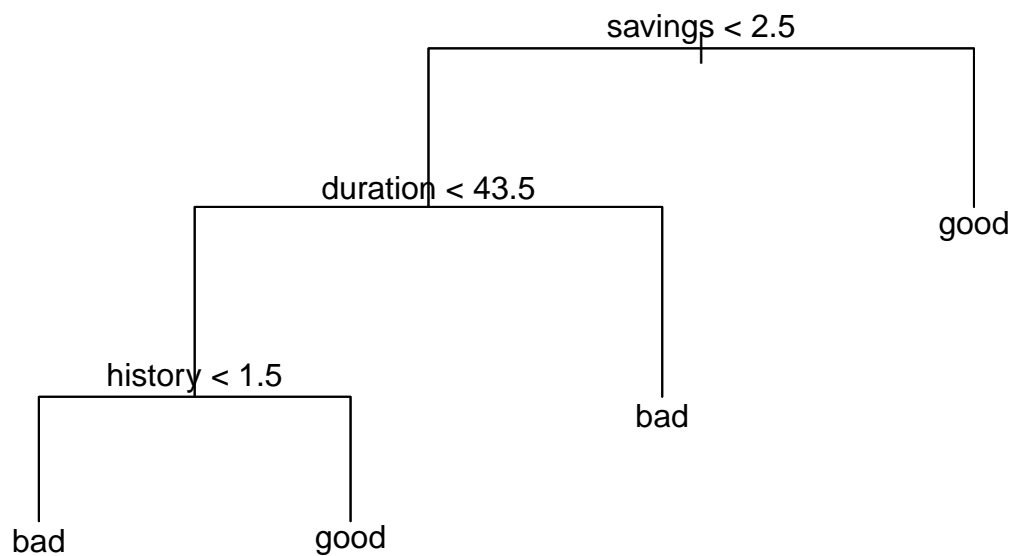
#> $confusion_matrix
#>      true
#> pred  bad good
#>   bad   59  32
#>   good  88 321
#>
#> $classification_rate
#> [1] 0.76

#> $confusion_matrix
#>      true
#> pred  bad good
#>   bad   18  26
#>   good  57 149
#>
#> $classification_rate
#> [1] 0.668
```

3



```
#> $confusion_matrix
#>      true
#> pred  bad good
#>   bad   22  12
#>   good  53 163
#>
#> $classification_rate
#> [1] 0.74
```



4

```

#> $confusion_matrix
#>      true
#> pred  bad good
#>   bad   95  98
#>   good  52 255
#>
#> $classification_rate
#> [1] 0.7

#> $confusion_matrix
#>      true
#> pred  bad good
#>   bad   50  61
#>   good  25 114
#>
#> $classification_rate
#> [1] 0.656

```

5

```

#> $confusion_matrix
#>      true

```



```

#> pred    bad good
#>   bad   95   98
#>   good  52  255
#>
#> $classification_rate
#> [1] 0.7

#> $confusion_matrix
#>      true
#> pred    bad good
#>   bad   50   61
#>   good  25  114
#>
#> $classification_rate
#> [1] 0.656

```

Appendix

Code for Assignment 1

```
library(ggplot2)
library(glmnet)

data <- read.csv("../data/australian-crabs.csv", sep=",")
ggplot(data) +
  geom_point(aes(x=RW, y=CL, color=sex))
LDA <- function(X, y) {
  n <- nrow(X)
  p <- ncol(X)

  labels <- unique(y)
  priors <- table(y) / length(y)

  means <- aggregate(X, list(y), mean)
  means <- as.matrix(means[, -1], ncol=p)

  lengths <- by(X, list(y), nrow)

  cov_mats <- by(X, list(y), cov)
  cov_mats <- lapply(1:length(lengths), function(i) {
    cov_mats[[i]] * lengths[[i]]
  })

  sigma <- as.matrix(Reduce("+", cov_mats) / sum(lengths), nrow=p)
  sigma_inv <- solve(sigma)

  w0 <- sapply(1:length(labels), function(i) {
    -(1 / 2) * t(means[i,]) %*% sigma_inv %*% means[i,] + log(priors[i])
  })

  w1 <- sapply(1:length(labels), function(i) {
    sigma_inv %*% means[i, ]
  })

  names(w0) <- levels(labels)
  colnames(w1) <- levels(labels)

  list(w0=w0, w1=w1, sigma=sigma)
}

X <- cbind(data$RW, data$CL)
y <- data$sex

result <- LDA(X, y)

w1 <- result$w1[, 2] - result$w1[, 1]
w0 <- result$w0[2] - result$w0[1]

intercept <- -w0 / w1[2]
```

```

slope <- -w1[1] / w1[2]
predicted <- as.numeric((w0 + w1 %*% t(X)) > 0)
predicted <- factor(predicted, levels=c(0, 1), labels=c("Female", "Male"))

plot_data <- data.frame(RW=data$RW, CL=data$CL, class=predicted)
line_data <- data.frame(intercept=intercept, slope=slope)

ggplot() +
  geom_point(data=plot_data, aes(x=RW, y=CL, color=class)) +
  geom_abline(data=line_data, intercept=intercept, slope=slope,
              color="black", linetype="dotted", size=1)
plot_data <- data.frame(RW=data$RW, CL=data$CL, class=data$sex)
line_data <- data.frame(intercept=intercept, slope=slope)

ggplot() +
  geom_point(data=plot_data, aes(x=RW, y=CL, color=class)) +
  geom_abline(data=line_data, intercept=intercept, slope=slope,
              color="black", linetype="dotted", size=1)
logistic_data <- data.frame(sex=as.numeric(data$sex) - 1, RW=data$RW, CL=data$CL)

glmfit <- glm(sex ~ RW + CL, data=logistic_data, family=binomial(link=logit))
coefficients <- coef(glmfit)

predicted <- as.numeric(glmfit$fitted.values > 0.5)
predicted <- factor(predicted, levels=c(0, 1), labels=c("Female", "Male"))

intercept <- -coefficients[1] / coefficients[3]
slope <- -coefficients[2] / coefficients[3]

plot_data <- data.frame(RW=data$RW, CL=data$CL, class=predicted)
line_data <- data.frame(intercept=intercept, slope=slope)

ggplot() +
  geom_point(data=plot_data, aes(x=RW, y=CL, color=class)) +
  geom_abline(data=line_data, intercept=intercept, slope=slope,
              color="black", linetype="dotted", size=1)

```

Code for Assignment 2

```

library(gdata)
library(tree)
library(partykit)
library(ggplot2)
library(reshape2)
library(e1071)

data_division <- function(n, training, test, validation) {
  indices <- 1:n

  train <- sample(indices, floor(n * 0.5))
  test <- sample(indices[-train], floor(n * test))
  validation <- indices[-c(train, test)]
}

```

```

    list(train=train, test=test, validation=validation)
}

data <- read.xls("../data/creditscoring.xls")
set.seed(12345)
indices <- data_division(nrow(data), 0.5, 0.25, 0.25)

train <- data[indices$train,]
test <- data[indices$test,]
validation <- data[indices$validation,]
prediction <- function(model, X, y) {
  predicted <- predict(model, X)

  if (is.matrix(predicted)) {
    predicted <- factor(ifelse(predicted[, 1] > predicted[, 2], 0, 1), levels=c(0, 1), labels=c("bad", "good"))
  }

  confusion_matrix <- table(pred=predicted, true=y)
  list(confusion_matrix=confusion_matrix,
        classification_rate=sum(diag(confusion_matrix)) / sum (confusion_matrix))
}

dtreefit <- tree(good_bad ~ ., data=train, split="deviance")
gtreefit <- tree(good_bad ~ ., data=train, split="gini")

prediction(dtreefit, train[, -ncol(train)], train$good_bad)
prediction(dtreefit, test[, -ncol(test)], test$good_bad)
prediction(gtreefit, train[, -ncol(train)], train$good_bad)
prediction(gtreefit, test[, -ncol(test)], test$good_bad)

leaves <- 2:summary(dtreefit)[4]$size
train_score <- rep(0, max(leaves))
validation_score <- rep(0, max(leaves))

for(i in leaves) {
  prunedTree <- prune.tree(dtreefit, best=i)
  pred <- predict(prunedTree, newdata=validation, type="tree")
  train_score[i] <- deviance(prunedTree)
  validation_score[i] <- deviance(pred)
}

plot_data <- data.frame(Leaves=leaves, Train=train_score[leaves], Validation=validation_score[leaves])
plot_data <- melt(plot_data, id="Leaves", value.name="Score", variable.name="Set")

ggplot() +
  geom_line(data=plot_data, aes(x=Leaves, y=Score, color=Set)) +
  scale_x_continuous(breaks=leaves)
optimal_leaves <- which.min(validation_score[leaves]) + 1
optimal_tree <- prune.tree(dtreefit, best=optimal_leaves)
prediction(optimal_tree, test[, -ncol(test)], test$good_bad)
plot(optimal_tree)
text(optimal_tree, pretty=0)

bayesfit <- naiveBayes(good_bad ~ ., data=train)

```

```
prediction(bayesfit, train[, -ncol(train)], train$good_bad)
prediction(bayesfit, test[, -ncol(test)], test$good_bad)

bayesfit <- naiveBayes(good_bad ~ ., data=train)
prediction(bayesfit, train[, -ncol(train)], train$good_bad)
prediction(bayesfit, test[, -ncol(test)], test$good_bad)
```