

Introduction to Machine Learning

Lab 3

Rasmus Holm


2016-11-24

Contents

Assignment 1	2
1	2
2	2
3	4
4	4
5	4
Assignment 2	5
1	5
2	7
3	8
4	10
Appendix	11
Code for Assignment 1	11
Code for Assignment 2	11


Assignment 1



1



The data looks like it could be modelled by a cubic spline function.

2





3


4

5


Assignment 2

1


```
#> [1] "95.385" "4.229" "0.186" "0.084" "0.072" "0.021" "0.007"
#> [8] "0.003" "0.003" "0.002" "0.001" "0.001" "0.001" "0.001"
#> [15] "0.000" "0.000" "0.000" "0.000" "0.000" "0.000" "0.000"
#> [22] "0.000" "0.000" "0.000" "0.000" "0.000" "0.000" "0.000"
#> [29] "0.000" "0.000" "0.000" "0.000" "0.000" "0.000" "0.000"
#> [36] "0.000" "0.000" "0.000" "0.000" "0.000" "0.000" "0.000"
#> [43] "0.000" "0.000" "0.000" "0.000" "0.000" "0.000" "0.000"
#> [50] "0.000" "0.000" "0.000" "0.000" "0.000" "0.000" "0.000"
#> [57] "0.000" "0.000" "0.000" "0.000" "0.000" "0.000" "0.000"
#> [64] "0.000" "0.000" "0.000" "0.000" "0.000" "0.000" "0.000"
#> [71] "0.000" "0.000" "0.000" "0.000" "0.000" "0.000" "0.000"
#> [78] "0.000" "0.000" "0.000" "0.000" "0.000" "0.000" "0.000"
#> [85] "0.000" "0.000" "0.000" "0.000" "0.000" "0.000" "0.000"
#> [92] "0.000" "0.000" "0.000" "0.000" "0.000" "0.000" "0.000"
#> [99] "0.000" "0.000" "0.000" "0.000" "0.000" "0.000" "0.000"
#> [106] "0.000" "0.000" "0.000" "0.000" "0.000" "0.000" "0.000"
#> [113] "0.000" "0.000" "0.000" "0.000" "0.000" "0.000" "0.000"
#> [120] "0.000" "0.000" "0.000" "0.000" "0.000" "0.000" "0.000"
#> [1] "0.954" "0.996" "0.998" "0.999" "1.000" "1.000" "1.000" "1.000"
#> [9] "1.000" "1.000" "1.000" "1.000" "1.000" "1.000" "1.000" "1.000"
#> [17] "1.000" "1.000" "1.000" "1.000" "1.000" "1.000" "1.000" "1.000"
#> [25] "1.000" "1.000" "1.000" "1.000" "1.000" "1.000" "1.000" "1.000"
#> [33] "1.000" "1.000" "1.000" "1.000" "1.000" "1.000" "1.000" "1.000"
#> [41] "1.000" "1.000" "1.000" "1.000" "1.000" "1.000" "1.000" "1.000"
#> [49] "1.000" "1.000" "1.000" "1.000" "1.000" "1.000" "1.000" "1.000"
#> [57] "1.000" "1.000" "1.000" "1.000" "1.000" "1.000" "1.000" "1.000"
#> [65] "1.000" "1.000" "1.000" "1.000" "1.000" "1.000" "1.000" "1.000"
#> [73] "1.000" "1.000" "1.000" "1.000" "1.000" "1.000" "1.000" "1.000"
#> [81] "1.000" "1.000" "1.000" "1.000" "1.000" "1.000" "1.000" "1.000"
#> [89] "1.000" "1.000" "1.000" "1.000" "1.000" "1.000" "1.000" "1.000"
#> [97] "1.000" "1.000" "1.000" "1.000" "1.000" "1.000" "1.000" "1.000"
#> [105] "1.000" "1.000" "1.000" "1.000" "1.000" "1.000" "1.000" "1.000"
#> [113] "1.000" "1.000" "1.000" "1.000" "1.000" "1.000" "1.000" "1.000"
#> [121] "1.000" "1.000" "1.000" "1.000" "1.000" "1.000" "1.000" "1.000"
```




2




3





4



Appendix

Code for Assignment 1

```
library(ggplot2)
library(tree)
library(reshape)

data <- read.csv2("../data/State.csv", header=TRUE, sep=";")
data <- data[order(data$MET),]
ggplot(data) +
  geom_point(aes(x=MET, y=EX))

tree.control(nobs=nrow(data), minsize=8)
treefit <- tree(EX ~ MET, data=data)

set.seed(12345)
treefit.cv <- cv.tree(treefit, FUN=prune.tree, K=10)
optimal_leaf_count <- treefit.cv$size[which.min(treefit.cv$dev)]

optimal_tree <- prune.tree(treefit, best=optimal_leaf_count)
plot(optimal_tree)
text(optimal_tree, pretty=0)
predicted <- predict(optimal_tree, data)
plot_data <- data.frame(MET=data$MET, Actual=data$EX, Estimate=predicted)
plot_data <- melt(plot_data, id.vars="MET")
names(plot_data) <- c("MET", "Data", "EX")

ggplot(plot_data) +
  geom_point(aes(x=MET, y=EX, color=Data))
residuals <- resid(optimal_tree)

plot_data <- data.frame(x=1:length(residuals), y=residuals)

ggplot(plot_data) +
  xlab("Index") +
  ylab("Residual") +
  geom_point(aes(x=x, y=y))
```

Code for Assignment 2

```
library(ggplot2)
library(fastICA)
library(pls)
library(reshape2)

data <- read.csv2("../data/NIRSpectra.csv")

X <- scale(data[, -ncol(data)])
y <- data[, ncol(data)]
```

```

traceplot <- function(n, pc1, pc2) {
  plot_data <- data.frame(x=1:n, PC1=pc1, PC2=pc2)
  plot_data <- melt(plot_data, id="x")
  names(plot_data) <- c("Index", "Component", "Value")
  xlimits <- seq(0, n, by=10)

  ggplot(plot_data) +
    geom_line(aes(x=Index, y=Value, color=Component), show.legend=FALSE) +
    scale_x_discrete(limits=xlimits) +
    facet_grid(Component ~ ., scales="free")
}

pca <- prcomp(X)

## Eigenvalues
lambda <- pca$sdev^2
variances <- lambda / sum(lambda)

var99_comp_count <- which.max(cumsum(variances * 100) > 99)
components <- as.data.frame(pca$x[, 1:var99_comp_count])

U <- pca$rotation
traceplot(nrow(U), U[, 1], U[, 2])

set.seed(12345)
ica <- fastICA(X, var99_comp_count, alg.typ = "parallel", fun = "logcosh", alpha = 1,
  method = "R", row.norm = FALSE, maxit = 200, tol = 1e-06, verbose = FALSE)

W_prime <- ica$K %*% ica$W
components <- as.data.frame(ica$S)
traceplot(nrow(W_prime), W_prime[, 1], W_prime[, 2])
ggplot(components) +
  geom_point(aes(x=V1, y=V2))

set.seed(12345)
pcrfit <- pcr(Viscosity ~ ., data=data, scale=TRUE)
cvpcrfit <- crossval(pcrfit, segments=10, segment.type="random")
cv_scores <- t(matrix(MSEP(cvpcrfit)$val, nrow=2))
plot_data <- data.frame(cbind(1:ncol(data), cv_scores))
colnames(plot_data) <- c("Components", "CV", "adjCV")
plot_data <- melt(plot_data, id="Components", variable_name="Measure")
names(plot_data)[ncol(plot_data)] <- "Score"
xlimits <- seq(0, ncol(data), by=5)

ggplot(plot_data) +
  geom_line(aes(x=Components, y=Score, color=Measure), size=1) +
  scale_x_discrete(limits=xlimits)

```